

Sravana Lakshmi Thota

Boston, MA | +1 617-318-7631 | Sraanalakshmithota3@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

AI Engineer with 2.5 years of production data engineering (Azure, PySpark, SQL) and graduate-level applied ML, focused on LLM-powered systems — multi-agent workflows, RAG, and LLM APIs. Build AI systems end to end, from data pipeline to deployed, self-monitoring model, in Python and the Claude and Gemini APIs.

TECHNICAL SKILLS

Languages: Python, SQL, R, JavaScript

AI / ML: PyTorch, TensorFlow, scikit-learn, Hugging Face, YOLOv8, RDKit, DECIMER, Computer Vision, Model Evaluation, RAG, LangChain, Vector Databases, LLM Integration (Claude, GPT), Anthropic SDK

Data Engineering: Azure Data Factory, Azure Databricks, PySpark, ETL Pipeline Design, SSIS

Cloud & Databases: Azure, AWS, Microsoft SQL Server, SSMS

APIs & Tools: FastAPI, Flask, REST APIs, Pydantic, Git, GitHub, Docling, Power BI, Tableau

PROFESSIONAL EXPERIENCE

KPMG Global Services Pvt. Ltd. — Bangalore, India

Aug 2022 – Dec 2024

Associate Consultant, Data Engineering

- Designed and deployed scalable ETL pipelines on Azure Data Factory, transforming multi-source enterprise data into analytics-ready datasets serving 10+ downstream reporting systems.
- Optimized complex SQL queries and stored procedures on Microsoft SQL Server, reducing execution time by up to 40% on critical reports.
- Automated data workflows with Azure Databricks and PySpark, improving pipeline reliability; awarded the SUPERTeam Award for high-quality data solutions.

Feminine Intelligence Agency (FIA) — Methodology Architect

Sep – Dec 2025

- Led evaluation methodology for a 35-member team building an AI-powered abuse-detection system; designed a counterfactual-testing framework and 5-dimensional rubric (Accuracy, Clarity, Tone, Safety, Respect) to improve LLM output robustness.

APPLIED ML PROJECTS

Multi-Agent AI Workflow System

github.com/sravanithota311/multi-agent-ai

- Built a multi-agent AI system using FastAPI with modular agents for retrieval, summarization, and action planning; integrated the Anthropic Claude API for LLM-driven decision support.
- Designed RESTful APIs and orchestration logic separating agent-level processing from workflow control; validated on enterprise use cases including construction delay analysis and resource optimization. Tech: Python, FastAPI, Anthropic Claude SDK, Pydantic, REST APIs.

Self-Distilling Classifier — LLM Distillation / MLOps

github.com/sravanithota311/self-distilling-classifier

- Built a self-retraining pipeline that uses the Gemini API as a labeling teacher to continuously train a lightweight local classifier on a live arXiv stream, reaching ~75–80% student-teacher agreement while cutting inference cost.
- Added champion/challenger promotion gating, data-drift monitoring, and a live observability dashboard (sravanithota311.github.io/self-distilling-classifier); fully automated on scheduled GitHub Actions. Tech: Python, scikit-learn, Gemini API, GitHub Actions.

OCSR Benchmark for Real-World Chemistry PDFs — Graduate Research Capstone

Apr – Jun 2026

- Built an end-to-end Python pipeline to benchmark DECIMER (transformer-based OCSR model) on 377 figures from 20 real-world medicinal chemistry journal papers, coordinating with a 6-member research team toward a co-authored publication.
- Implemented RDKit-based precision/recall evaluation against hand-verified ground truth; confirmed a systematic 0% recall failure mode across three OCSR tools and categorized failure modes (hallucinations, multi-molecule concatenation, 41% unparseable rate). Tech: Python, DECIMER, RDKit, PyTorch, Docling, Pandas.

EDUCATION

Northeastern University — Boston, MA

Jan 2025 – Jun 2026

Master of Professional Studies in Analytics (Applied Machine Intelligence), GPA: 3.86/4.0

CERTIFICATIONS

Microsoft Certified: Azure Fundamentals (AZ-900); BCG X GenAI & Vista AI in Action Job Simulations (Forage); IBM AI Fundamentals; Prompt Engineering for GenAI (LinkedIn).